

Java Thread Synchronization

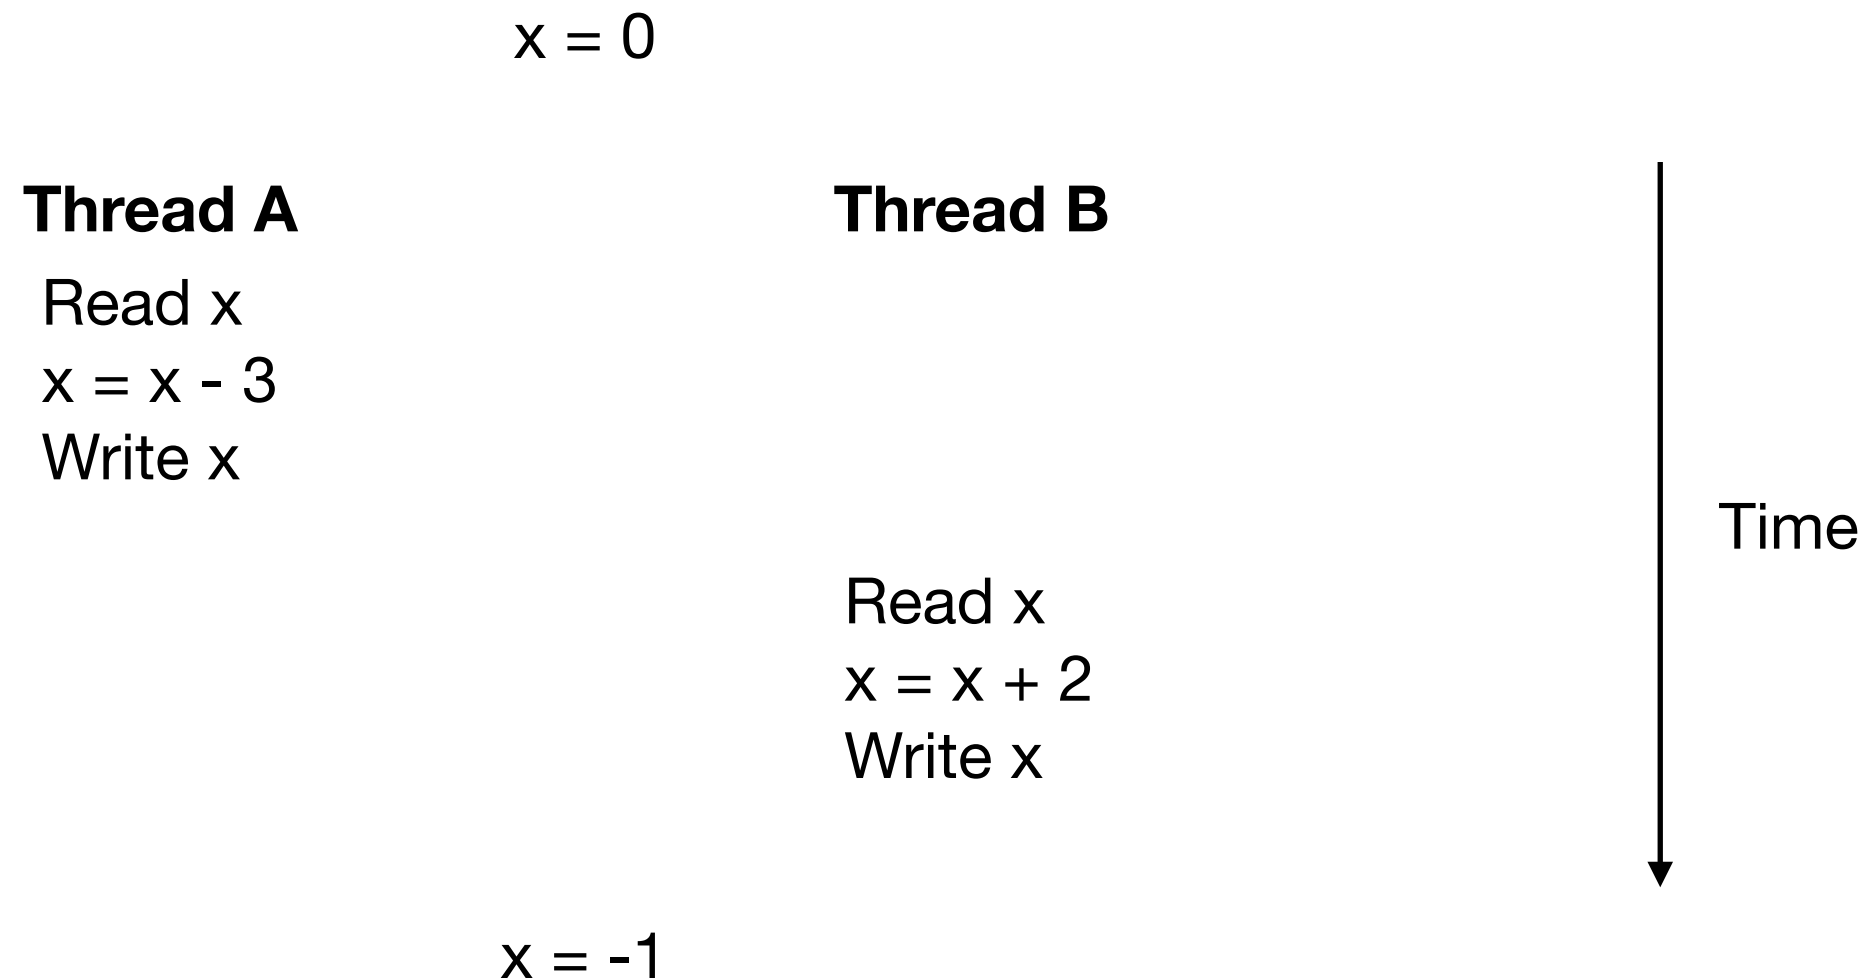
Terminology

- Sequential (serial) computing
 - One thread, instructions executed one at a time
- Concurrent computing
 - Multiple processes (or threads)
 - Instruction execution overlaps but is not simultaneous
 - Improves throughput (from system perspective) over sequential computing
- Parallel computing
 - Simultaneous execution
 - Ex threads on multi-core node or processes on multiple nodes

Synchronization

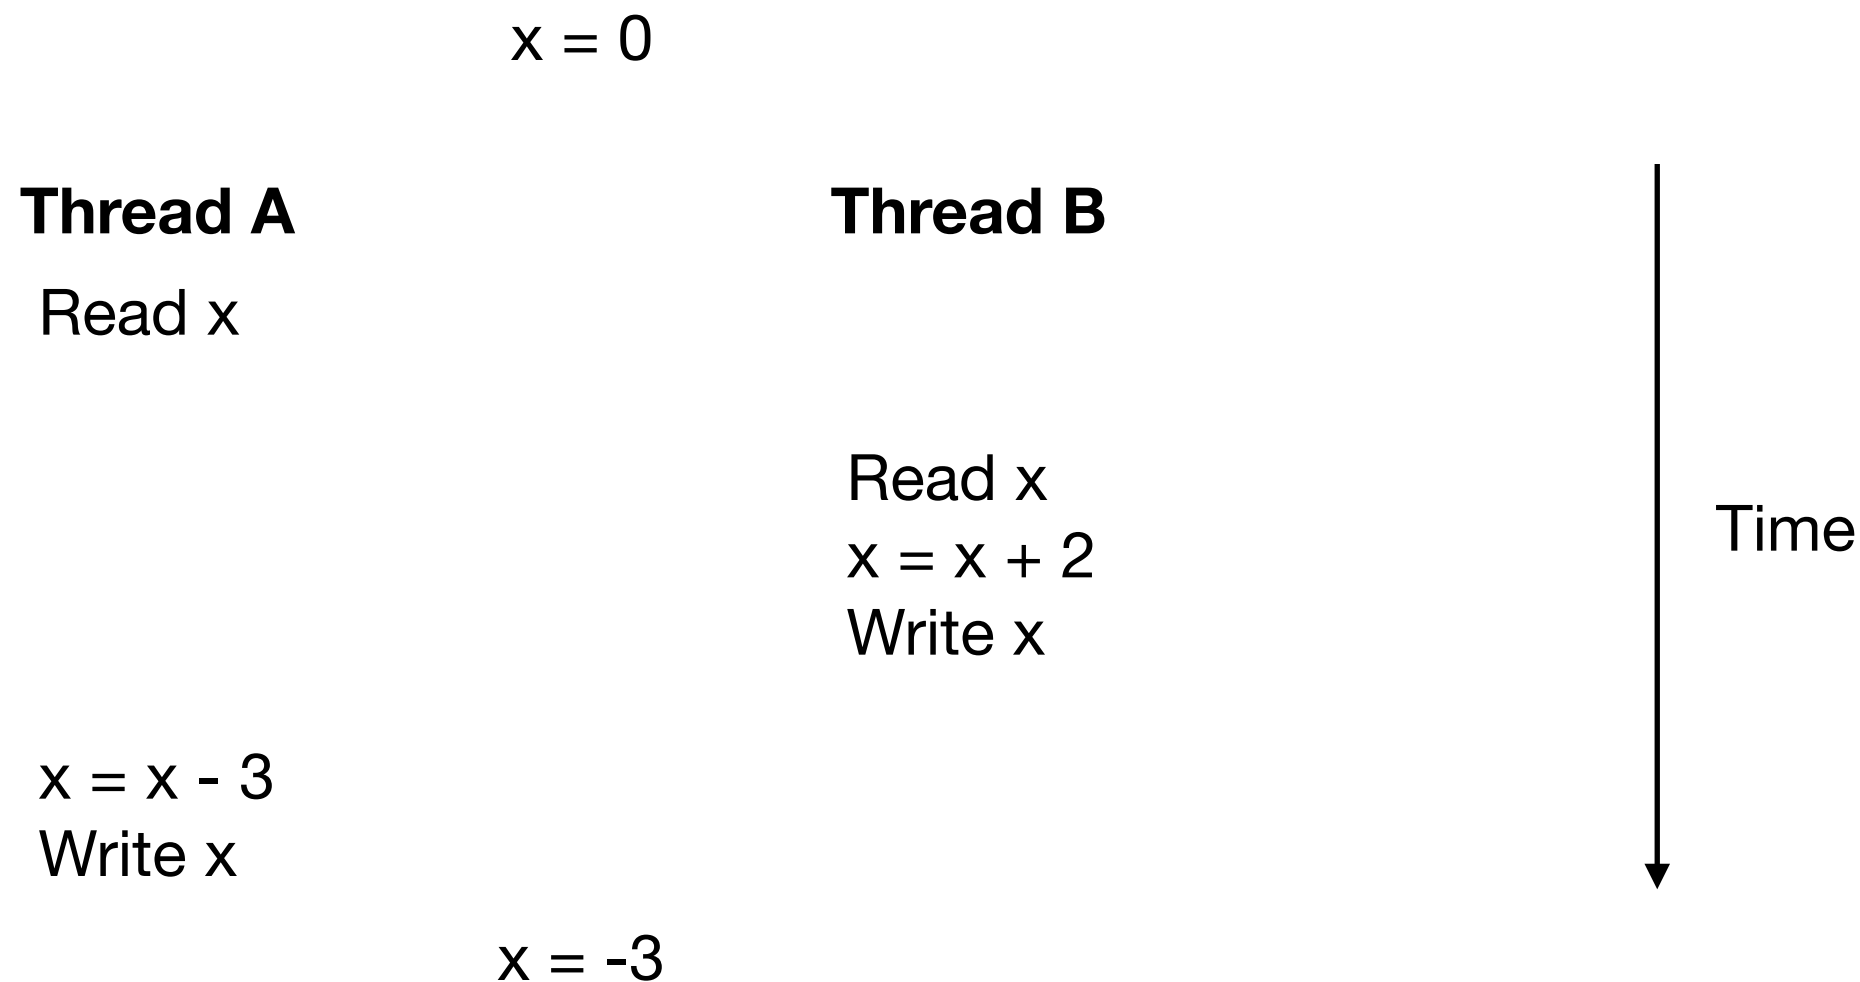
- Processes (and threads)
 - Coordinating activities to achieve some goal
- Data
 - Avoiding unintended modification (corruption) of data

No Problems



These threads are synchronized or read/modify/write is atomic.

Lost Update Problem



The $x=x+2$ update is lost.

Inter-leaving threads can lead to race conditions (the behavior (output) depends on a sequence or timing of events).

Critical Section

- A part of a program where concurrent access may result in unexpected behavior or non-deterministic results
 - Ex. A section of code that accesses a data structure or a hardware device
- Need to take steps to enforce mutually exclusive access
 - Avoid race conditions
- Programmers should design critical sections to be as short as possible.

Dining Philosophers Problem

- Classic problem illustrates synchronization issue presented by Edsger Dijkstra (1965)
- 5 silent philosophers sit around a table with 5 forks between them (1 fork between each pair of philosophers)
- Philosophers think or eat. Eating requires two forks.
- Must put down both forks after eating.
- Find an algorithm so that no philosopher starves.

Types of Locks

- A lock is a synchronization mechanism that enforces mutual exclusion for a critical section
- Semaphore
 - Invented by Edsger Dijkstra (early 1960s)
 - A variable or object that keeps track of how many of a resource are in use
- Binary Semaphore
 - A semaphore initialized to 1
- Monitor

Java Support for Mutual Exclusion

- `synchronized` statement
- `synchronized` method modifier
- `java.util.Collections` (some methods)
- `java.util.concurrent.Semaphore` class

synchronized statement

- Syntax: `synchronized (expression) statement`
 - *expression* must resolve to an object or an array
 - *statement* is the code of the critical section
- JVM does not execute critical section until it obtains an exclusive lock on *expression*
- JVM maintains exclusive lock until critical section is completed

synchronized example

```
public static void SortIntArray(int[] a) {  
    synchronized(a) {  
        // do the array sort here  
    }  
}
```

- This is synchronized so that some other thread can't change elements of the array during the sort.
- At least not other threads that protect their changes to the array while synchronized.

Thread-safe Example

```
public class MyIntList {  
    private int[] list;  
  
    public MyIntList(int size) {  
        list = new int[size];  
    }  
  
    public void sort() {  
        synchronized(list) {  
            // do the sort  
        }  
    }  
  
    public void set(int index, int val) {  
        synchronized(list) {  
            list[index] = val;  
        }  
    }  
}
```

```
    public int get(int index) {  
        synchronized(list) {  
            return list[index];  
        }  
    }  
  
    public int size() {  
        return list.length;  
    }  
}
```

synchronized method modifier

- The entire method is a critical section
- JVM obtains an exclusive lock
 - synchronized **static** methods obtain lock on **class**
 - synchronized **instance** methods obtain lock on **object**
- No other thread can execute any synchronized method
- Any other thread can execute non-synchronized methods

Still Thread-safe

```
public class MyIntList {  
    private int[] list;  
  
    public MyIntList(int size) {  
        list = new int[size];  
    }  
  
    public synchronized void sort() {  
        // do the sort  
    }  
  
    public synchronized void set(int index, int value) {  
        list[index] = value;  
    }  
  
    public synchronized int get(int index) {  
        return list[index];  
    }  
  
    public int size() {  
        return list.length;  
    }  
}
```

java.util.concurrent.Semaphore

- A counting semaphore. For controlling access to a limited pool of resources
- Example. Given a pool of 5 resources and 10 threads, a counting semaphore ensures that no more than 5 resources are allocated at a time. If the pool is fully allocated, the next thread to request a resource is blocked until another thread returns its resource.
- Semaphore.acquire() if count > 0, count = count - 1
 else block
- Semaphore.release() count++