

# Getting Started with Automatic Compiler Vectorization

# Parallelism is Key to Performance

- Types of parallelism
  - Task-based (MPI)
  - Threads (OpenMP, pthreads)
  - Vector processing

# Scalar Processing

- Processor uses one element from each array on each loop iteration.

```
#define N 32
```

```
float a[N], b[N], c[N];
```

```
for (int i = 0; i < N; i++) {  
    a[i] = b[i] + c[i];  
}
```

Registers

b[i]

c[i]

a[i]

# Vector (SIMD) Processing

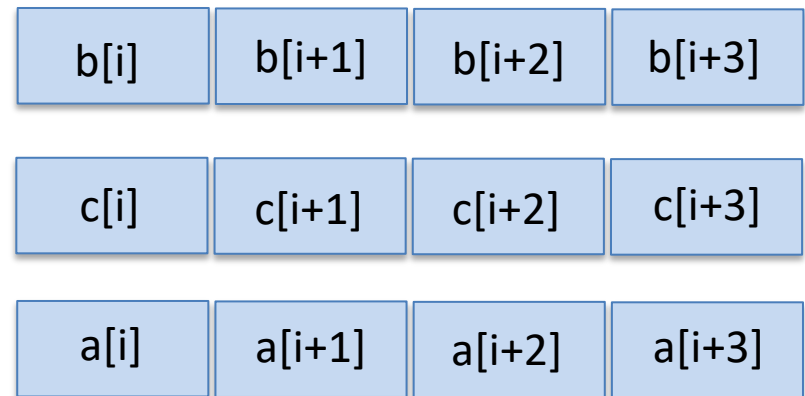
- Compiler unrolls the loop body a number of times.

```
#define N 32
```

```
float a[N], b[N], c[N];
```

```
for (int i = 0; i < N; i=i+4) {  
    a[i] = b[i] + c[i];  
    a[i+1] = b[i+1] + c[i+1];  
    a[i+2] = b[i+2] + c[i+2];  
    a[i+3] = b[i+3] + c[i+3];  
}
```

Registers

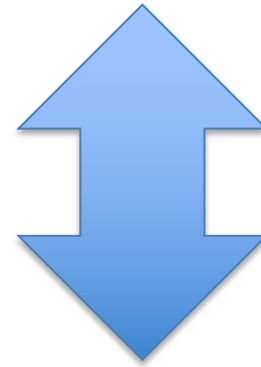


- Processor performs the same operation on a number of scalars (simultaneously)

# Techniques

Ease of Use / Portability

- Automatic Compiler Vectorization
- Compiler hints (directives / #pragmas)
- SIMD / Vector intrinsics
- Assembler instructions



Programmer Control

# Requirements / Guidelines

- Compiler optimization –O2 or higher
- Loop trip count known at runtime.
- A single entry and single exit
- Straight-line code
- Inner-most loop
- No function calls

# Manage the Compiler Report

- Vectorization is enable by default
- -qopt-report-phase=vec
- -qopt-report=[1-5]
  - Provides increasing levels of detail
- -qopt-report-routine:func1, func2

# Exercise

Test.cpp

```
#include <math.h>
void foo(float * theta, float * sth) {
    for (int i = 0; i < 128; i++) {
        sth[i] = sin(theta[i] + 3.1415927);
    }
}
```

```
$ icpc -qopt-report-phase=vec \
    -qopt-report=5 \
    -qopt-report-routine:foo \
    -c \
    Test.cpp
```



# Test.opttrpt (1)

Begin optimization report for: foo(float \*, float \*)

Report from: Vector optimizations [vec]

LOOP BEGIN at Test.cpp(3,5)

<Multiversed v1>

...

remark #15300: LOOP WAS VECTORIZED

...

LOOP END

Good

LOOP BEGIN at Test.cpp(3,5)

<Multiversed v2>

remark #15304: loop was not vectorized: non-vectorizable loop instance from multiversioning

LOOP END

Bad

# Aliasing

- Add a command-line option to tell compiler that function arguments are not aliased.

```
$ icpc -qopt-report-phase=vec \  
    -qopt-report=5 \  
    -qopt-report-routine:foo \  
    -c \  
    -fargument-noalias \  
    Test.cpp
```

- Results in a single vectorized version of loop.

# Test.opttrpt (2)

Begin optimization report for: foo(float \*, float \*)

Report from: Vector optimizations [vec]

LOOP BEGIN at T1.cpp(3,5)

remark #15389: vectorization support: reference theta has unaligned access [ T1.cpp(4,18) ]

remark #15389: vectorization support: reference sth has unaligned access [ T1.cpp(4,9) ]

remark #15381: vectorization support: unaligned access used inside loop body

remark #15305: vectorization support: vector length 4

remark #15309: vectorization support: normalized vectorization overhead 0.081

remark #15417: vectorization support: number of FP up converts: single precision to double precision 1 [ T1.cpp(4,18) ]

remark #15418: vectorization support: number of FP down converts: double precision to single precision 1 [ T1.cpp(4,9) ]

remark #15300: LOOP WAS VECTORIZED

remark #15450: unmasked unaligned unit stride loads: 1

remark #15451: unmasked unaligned unit stride stores: 1

remark #15475: --- begin vector loop cost summary ---

remark #15488: --- end vector loop cost summary ---

LOOP END

# Alignment

- Tell the compiler that arrays are aligned.

```
#include <math.h>
void foo(float *theta, float *sth) {
    __assume_aligned(theta, 32);
    __assume_aligned(sth, 32);
    for (int i = 0; i < 128; i++) {
        sth[i] = sin(theta[i] + 3.1415927);
    }
}
```

# Test.optprt (3)

LOOP BEGIN at Test.cpp(3,5)

remark #15388: vectorization support: reference theta has aligned access [ Test.cpp(6,18) ]

remark #15388: vectorization support: reference sth has aligned access [ Test.cpp(6,9) ]

remark #15305: vectorization support: vector length 8

remark #15309: vectorization support: normalized vectorization overhead 0.006

remark #15417: vectorization support: number of FP up converts: single precision to double precision 1 [ Test.cpp(6,18) ]

remark #15418: vectorization support: number of FP down converts: double precision to single precision 1 [ Test.cpp(6,9) ]

remark #15300: LOOP WAS VECTORIZED

remark #15448: unmasked aligned unit stride loads: 1

remark #15449: unmasked aligned unit stride stores: 1

remark #15475: --- begin vector loop cost summary ---

remark #15476: scalar loop cost: 113

remark #15477: vector loop cost: 20.000

remark #15478: estimated potential speedup: 5.640

remark #15482: vectorized math library calls: 1

remark #15487: type converts: 2

remark #15488: --- end vector loop cost summary ---

LOOP END

# Avoid type conversions

- Specify float (f) on function call and literal value.

```
#include <math.h>
void foo(float *theta, float *sth) {
    __assume_aligned(theta, 32);
    __assume_aligned(sth, 32);
    for (int i = 0; i < 128; i++) {
        sth[i] = sinf(theta[i] + 3.1415927f);
    }
}
```

# T1.optprt (4)

LOOP BEGIN at Test.cpp(5,5)

remark #15388: vectorization support: reference theta has aligned access [ Test.cpp(6,18) ]

remark #15388: vectorization support: reference sth has aligned access [ Test.cpp(6,9) ]

remark #15305: vectorization support: vector length 8

remark #15309: vectorization support: normalized vectorization overhead 0.013

remark #15300: LOOP WAS VECTORIZED

remark #15448: unmasked aligned unit stride loads: 1

remark #15449: unmasked aligned unit stride stores: 1

remark #15475: --- begin vector loop cost summary ---

remark #15476: scalar loop cost: 110

remark #15477: vector loop cost: 9.870

remark #15478: estimated potential speedup: 11.130

remark #15482: vectorized math library calls: 1

remark #15488: --- end vector loop cost summary ---

LOOP END

# Resources

- <https://software.intel.com/en-us/articles/vectorization-diagnostics-for-intelr-c-compiler-150-and-above>
- <https://software.intel.com/en-us/articles/vectorization-essential>
- <https://software.intel.com/sites/default/files/8c/a9/CompilerAutovectorizationGuide.pdf>
- <https://software.intel.com/en-us/articles/get-a-helping-hand-from-the-vectorization-advisor>
- <https://software.intel.com/en-us/mkl>
- <https://software.intel.com/en-us/intel-parallel-universe-magazine>